Ralph - A Dylan dialect that compiles to JavaScript

Bastian Mueller
IT University of Copenhagen, Denmark
bmue@itu.dk

ABSTRACT

We present the object-centered dynamic functional programming language Ralph, which is inspired by Lisp. Ralph supports an extended subset of Dylan's features (the intermediate Dylan standard with a prefix syntax) and compiles to JavaScript. The Ralph compiler produces more efficient JavaScript code than similar compilers translating Lisp-like languages due to some trade-offs and utilization of an intermediate representation based on the Administrative normal form. We discuss the Ralph language, its mechanics, implementation of its features in JavaScript, and its multi-pass compilation approach.

1. INTRODUCTION

Ralph[1] is a programming language that targets JavaScript, more specifically its subset ECMAScript [6]. Ralph is a Lisp dialect mainly inspired by Dylan [13] as specified in the first edition of its manual [12] (from here on referred to as "early specification"). Dylan is based on Scheme, Common Lisp, and Smalltalk. Like Scheme and unlike Common Lisp, Dylan has a single namespace for functions and variables. The object system is derived from the Common Lisp Object System (CLOS), featuring multiple inheritance and generic functions performing multiple dispatch.

Dylan was invented by Apple and its original code-name was "Ralph". The same name was chosen for the new language to imply the effort of a partial reimplementation. Unlike the final and standardized version of Dylan which features an ALGOL-like syntax, this early version still used a syntax with prefix notation based on s-expressions [11].

Ralph provides many of Dylan's features, special forms, macros, and functions of the standard library. However, it has a simplified abstraction for collections, condition system and numerical tower. It provides a module system, procedural hygienic macros and easy interoperability with

---

[1]Available at `https://github.com/turbolent/ralph`

JavaScript. The class-based object system currently supports single inheritance and single dispatch.

Today it is possible to write rich applications for both desktop and mobile devices using Web technologies. Modern browsers offer a number of APIs for this purpose, but JavaScript is the only programming language able to use them. Besides the usage on the client-side, JavaScript has also become popular for server-side applications, through environments such as Node.js. JavaScript execution environments and their underlying engines have evolved from simple interpreters to sophisticated virtual machines involving Just-In-Time emission of native code.

JavaScript is a functional programming language influenced by Scheme and Self: it supports first-class functions, closures, lexical scoping, higher-order programming, dynamic typing and prototype based object-oriented programming.

Because of these features, an increasing number of compilers are targeting JavaScript, both for existing languages, new languages or low-level byte-code. For example, the Google Web Toolkit allows writing applications in Java, Parenscript is a compiler for a subset of Common Lisp, and ClojureScript is a compiler for a subset of Clojure. CoffeeScript provides an alternative syntax for JavaScript. Emscripten is a compiler translating LLVM bytecode to JavaScript.

Ralph was designed for use in application development. It is dynamic to allow rapid prototyping and development, and has a focus on expressiveness and efficiency. Much like Dylan, it tries to bring the advanced features of Lisp to non-Lisp programmers, having a good trade-off between sophisticated features and performance. Performance and efficiency are especially crucial on mobile devices, which have only limited processing capabilities and reducing power consumption is important. The compiler of Ralph generates safe and efficient JavaScript close to hand-written code, due to trade-offs, utilization of an adequate intermediate representation and high-level optimizations.

We motivate the problem set by presenting the issues associated with using JavaScript in Section 2, explaining why it is desirable to have a higher-level language targeting JavaScript. Section 3 is looking into related projects and compilation approaches. Afterwards we describe our solutions for several problems in Section 4. We present Ralph's limitations in Section 5 and describe future work and conclude in Section 6.

2. PROBLEMS WITH JAVASCRIPT

JavaScript has several problematic features which programmers constantly have to avoid [5]. A lexical environment is only created for functions, `with` statements and `catch` clauses of `try` statements. Block scope or an equivalent of a let-expression are not available. Another common problem is hoisting. Variable and function declarations (but not their initializers) are moved to the top of their enclosing scope and default to `undefined`. An example is shown in Listing 1.

```
1  function foo() {
2      bar();
3      var x = 1;
4  }
```
(a) source

```
1  var foo;
2  foo = function () {
3      var x;
4      bar();
5      x = 1;
6  }
```
(b) hoisted

Listing 1: Hoisting of declarations

The problematic consequences of these features is illustrated in the following two programs. In the example shown in Listing 2, it could be assumed the conditional in `bar` fails, due to the initialization of `foo` to 1. However, because of hoisting, the conditional succeeds.

```
1  var foo = 1;
2  function bar() {
3      if (!foo) {
4          var foo = 10;
5      }
6      return foo;
7  }
8  bar();  // ⇒ 10
```
Listing 2: Bug caused by hoisting

In the second example shown in Listing 3, it could be assumed the assignment in line 3 would affect the global variable, but the function declaration is hoisted and a local variable `a` is introduced inside function `b`, even though the declaration of function `a` is actually never reached.

```
1  var a = 1;
2  function b() {
3      a = 10;
4      return;
5      function a() {}
6  }
7  b();
8  a // ⇒ 1
```
Listing 3: Hoisting with unreachable code

`For` loops are not creating new bindings in each iteration. In combination with closures this leads to another common problem that is exemplified in Listing 4.

```
1  var fns = [];
2  for (var i = 0; i < 5; i++) {
3    fns.push(function () {
4      return i;
5    });
6  }
7  fns[3]();  // ⇒ 5
```
Listing 4: For loop in connection with closure

The initialization part of the `for` loop is hoisted and in the step part, the variable is incremented. The closure added to the array is containing a reference to the variable, not the value that was bound when the closure was created.

Another common problem is JavaScript's uncommon semantics of truth and comparison using the equality operator `==`, which involves several coercion rules if its operands are of different types. Some examples demonstrating the behavior are shown in Listing 5.

```
1  '' == '0'          // ⇒ false
2  '' == 0            // ⇒ true
3  '0' == 0           // ⇒ true
4  false == '0'       // ⇒ true
5  ' \t\n ' == 0      // ⇒ true
6  ' \n ' ? true : false  // ⇒ true
```
Listing 5: Examples of equality and truth behavior

The language is also missing a mechanism to organize code and create reusable software components that safely interact with each other. For example, Common Lisp provides packages for this purpose. Dylan provides modules and libraries. In JavaScript, all scripts loaded into an environment are evaluated in the global scope, thus leading to possible interference between programs using identical names.

Even though JavaScript has problems and writing programs directly is error-prone, using JavaScript as a compilation target is viable.

3. RELATED WORK

Besides Ralph, more compilers target JavaScript, both for full or substantial subsets of existing languages and new languages[2]. We describe a Lisp dialect, thus this section will survey some other Lisp dialects targeting JavaScript[3].

In general, compilers follow either one of two approaches. The first and most common is compiling the high-level source language to high-level JavaScript. For example, Parenscript, ClojureScript and scheme2js are in this category. Another approach is using an existing compiler for the source language and compiling the emitted low-level language to JavaScript. For instance, Gambit-JS is a compiler for Scheme relying on the Gambit-C compiler for this purpose.

Parenscript is a compiler for an extended subset of Common Lisp and is available as a Common Lisp library. Its goals are efficiency and good interoperability with JavaScript. Generated code is readable and has no run-time dependencies, but the standard library is small. Programs written in Parenscript can also be used as Common Lisp without modification. Even though some optimizations are applied, it is producing code naively in a direct-style manner.

ClojureScript is a compiler for a subset of Clojure and is written in Clojure itself. It only employs few optimizations and also produces JavaScript naively. It uses the run-time

---

[2]For a comprehensive list, see http://altjs.org/
[3]A feature matrix of more Lisp-like languages compiling to JavaScript is available at http://ceaude.twoticketsplease.de/js-lisps.html

of the Google Web Toolkit, and the compiler annotates the generated code with type hints in comments. The resulting code is further compiled by the Google Closure Compiler. It performs a JavaScript-to-JavaScript translation, optimizing and minimizing the code as much as possible. For this purpose it is using the provided type hints and certain code conventions. However, it does not always generate efficient code, since higher-level information about the intentions of the programmer are lost during the translation. Another disadvantage is the large run-time dependency. Both Parenscript and ClojureScript allow the usage of macros and the latter is even hygienic, but they need to be written in the host language.

Other implementations in this category have similar characteristics and may be written in JavaScript instead. As ClojureScript and Parenscript, they commonly provide easy means of interoperating with native JavaScript code, by using built-in types and using the native calling convention. They commonly also generate readable code comparable to hand-written one, making it easier to debug. However, the majority of them is only performing few optimizations and naively generating JavaScript, resulting in modest performance. One particular problem is the generation of too many unnecessary closures, which has a large performance impact. To implement block-level scope, these compilers produce an anonymous function wrapper that is immediately invoked, as shown in Listing 6.

```
1  (setf x
2      (let ((x 1))
3        x))
4  ────────────────────
5  x = (function () {
6       var x = 1;
7       return x;
8     })();
```
(a) Parenscript

```
1  (let [x (let [x 1]
2               x)]
3    x)
4  ────────────────────
5  var x__2 =
6    (function () {
7      var x__1 = 1;
8      return x__1;
9    })();
```
(b) ClojureScript

Listing 6: Generation of closures for nested bindings

These closures are also generated for forms in expressions for which JavaScript statements need to be generated. An example for ClojureScript is shown in Listing 7. A binding for the result of an exception handling form is introduced. As the compiler needs to emit a JavaScript `try`/`catch` statement, it needs to be wrapped in an immediately invoked anonymous function to be used as an expression.

```
1  (let [x (try ...
2             (catch ...))]
3    ...
4  ────────────────────
5  var x__1 = (function () {
6      try { ... }
7      catch (...) { ... }
8    })();
9  ...
```
Listing 7: Generation of closures for forms resulting in JavaScript statements

A more sophisticated compiler for Scheme is scheme2js. It supports many features of R^5RS [1], but has a focus on efficiency rather than conformance. It performs tail-call optimization by translation to iterative constructs and using trampolines. It provides a non-hygienic macro system and supports first-class continuations (`call/cc`), implemented through exceptions. The special Replay-C [10] algorithm was developed for this optimization. Using exceptions in generated code has performance implications though, which are discussed in Section 5. Not relying on built-in types also results in an overhead when interoperating with native JavaScript code.

An example for the second approach, compiling low-level code to JavaScript, is Gambit-JS [15]. First, Gambit-C is used to compile Scheme code conforming to R^5RS to a final control flow graph (CFG), called the Gambit Virtual Machine. This intermediate representation (IR) is normally further compiled to C. Gambit-JS is emulating the low-level semantics of the instructions by implementing a custom stack in JavaScript. This approach has many advantages. Continuations are serializable and can be migrated between environments, proper tail-calls are implemented without resorting to trampolines, and generated code is faster than that of scheme2js. Major disadvantages are its complexity and difficulty of implementation, as well as more difficult interoperability.

4. SOLUTION
In this Section we describe the implementation details of Ralph and the design considerations for Ralph's implementation, taking into account the previous sections. We focus on JavaScript interoperability, and describe modules, the object system, macros and scoping. Finally we present the separate passes of the compiler and implemented optimizations.

4.1 Interoperability
For a language compiling to JavaScript interoperability is crucial, because execution environments such as the browser provide a number of APIs. Additionally third-party libraries can be called easily.

In Ralph we rely on JavaScript's native data structures and calling convention for performance reasons. Neither automatic coercion, which imposes a performance penalty, nor error-prone manual coercion is needed when performing calls between Ralph and JavaScript.

The root of the class hierarchy `<object>` is represented by the JavaScript type `Object`. It provides a mechanism to associate strings ("properties") with values of any type. For symbols, a class with the properties `name` and `module` is used.

Arrays are the fundamental data structure to represent collections in JavaScript. They dynamically grow when adding or removing elements and are efficiently implemented in the execution environments. Ralph's lists are represented by JavaScript arrays, instead of implementing a custom `Pair` type with `head` and `tail` properties, which would result in higher memory consumption and worse performance.

Numbers and boolean values are directly represented by their JavaScript counterparts. There is no explicit distinction between integer and floating point numbers, in contrast to Dylan's numerical tower.

Strings are represented by its respective JavaScript type as well, for the same reasons of speed and interoperability. The drawback is immutability, compared to having mutable strings as standardized in Dylan.

4.2 Module system

Dylan allows structuring code into reusable software component by means of modules and libraries. A library consists of one or more modules that contain the actual definitions. Each module acts as an independent namespace for identifiers. In the early specification, no module system was specified. We chose to adopt modules to group definitions. Each file is a module with its own namespace for identifiers. Definitions can be imported from other modules and exported using the `define-module` form. An example is shown in Listing 8.

```
1  (define-module ralph/reader
2    import: (ralph/stream
3            (ralph/format
4              only: (format-to-string))
5            (ralph/regexp
6              rename: (match match-regexp)))
7    export: (read))
```

**Listing 8:** Definition of the reader module

4.3 Object System

4.3.1 Classes

In comparison to Dylan, which provides a multiple inheritance class-based object system, Ralph currently provides single inheritance. The reason for this design decision is performance.

We previously implemented multiple inheritance without using the prototypial features of the object system provided by JavaScript. This turned out to be impractical in real-world applications due to poor performance. In the future we want to extend Ralph with multiple inheritance.

To understand how Ralph's object system is implemented, first the underlying prototype-based system is explained briefly. In JavaScript, every object might have a `__proto__` property referencing another object, called the prototype. When a property is accessed on an object which does not contain the property, and the object has a prototype, this is used instead for the lookup. For example, ({__proto__: {x: 1}}).x performs the lookup of property `x` on an object with prototype {x: 1}.

Functions are first-class objects and constructors are implemented as functions. Using the `new` operator, a fresh object is instantiated. The execution environment allocates a new object and sets its prototype to the constructor's `prototype` property, which is initially an empty object. If a property lookup on the instance fails, the value of the constructor's `prototype` property will be used instead. When a constructor is called using `new`, the implicit variable `this` refers to the instance is bound inside the body. This is also the case when

a function is called as a method, i.e., using property lookup in form of `object.function()`, rather than `function(object)`. This mechanism is illustrated in Listing 9.

```
1  function Person(name) {
2    this.name = name;
3  }
4  Person.prototype.hello = function () {
5    return "Hello, " + this.name + "!";
6  }
7  var john = new Person("John");
8  // john.__proto__ = Person.prototype
9  john.hello() // ⇒ "Hello, John!"
```

Listing 9: Using prototypes in JavaScript

This already resembles typical class-based object orientation when assuming the constructor as a "class". Single inheritance can be implemented using
`class.prototype.__proto__ = superclass.prototype`

In Ralph, the macro `define-class` produces a constructor. Additionally an object containing all properties (cf. slots in Dylan and CLOS) of the class is generated. Finally, a call to the runtime function `%make-class` is generated and these values and the superclass are passed as arguments. This is demonstrated in Listing 10.

```
1  (define-class <person> (<object>)
2    name)
3
4  (define <person>
5         (%make-class <object>
6                      (%object "name" #f)
7                      (%method <person> ())))
```

Listing 10: Code produced by macro `define-class`

At runtime, `%make-class` will setup the inheritance relationship by setting the prototype property as shown above, and keeping a reference to the superclass. Also, the prototype of the object containing all properties is set to that of the superclass, so that looking up properties of the class will also return inherited ones, most specific first. The definitions are shown in Listing 11.

```
1  (define-function %inherit (class superclass)
2    (set! (get class "%superclass")
3          superclass)
4    (set! (get class "prototype" "__proto__")
5          (get superclass "prototype"))
6    (set! (get class "%properties" "__proto__")
7          (get superclass "%properties"))
8    class)
9
10  (define-function %make-class
11      (superclass properties constructor)
12    (set! (get constructor "%properties")
13          properties)
14    (when superclass
15      (%inherit constructor superclass))
16    constructor)
```

Listing 11: Definitions of `%make-class` and `%inherit`

Similar to Dylan, the method `make` will allocate an instance of the passed class, and `initialize` will perform the actual initialization of all properties. If no values are passed to

`make`, the default forms, stored as functions, are used instead. Listing 12 contains the definitions of both functions.

```
1  (define-function make (type #rest arguments)
2    (bind ((object (%native-call "new" type)))
3      (apply initialize object arguments)
4      object))
5
6  (define-function %initialize-property
7      (object properties arguments property)
8    (unless (has? (get <object> "prototype")
9                  property)
10     (bind ((value
11            (or (get arguments property)
12                (bind ((value
13                        (get properties
14                             property)))
15                  (when value
16                    (value))))))
17       (set! (get object property) value))))
18
19 (define-method initialize
20     ((object <object>) #rest rest)
21   (bind ((arguments (as-object rest)))
22     (if-bind (properties
23               (get (type object)
24                    "%properties"))
25       (do (curry %initialize-property
26                  object properties
27                  arguments)
28         (object-properties
29          properties inherited?: #t))))
30    object)
```

Listing 12: Definitions of `make` and `initialize`

Recently, means of defining properties were added to JavaScript, including default values and access restrictions, using `Object.defineProperty`. The approach described above was chosen, as this feature is not yet available in all execution environments and its performance is poor compared to normal functions which set and get properties.

4.3.2 Functions

There are two types of functions: methods and generic functions. A generic function can contain one or more methods. Both generic functions and methods consist of a typed argument list and code. Similar to CLOS, functions are first-class citizens, and not defined inside of classes. When a method is called, its body is directly executed. When a generic function is called, the first argument is compared to the typed argument lists of all defined methods of the generic function and the most specific one is invoked. This is the means of object-oriented (single) dispatch.

Dylan performs multiple dispatch, i.e., all arguments and all elements of the argument lists are considered when determining the applicable method. Ralph is currently limited to single dispatch, i.e., only the first argument and first type of each argument list is taken into consideration.

Similar to the implementations of the class hierarchy, this design choice is due to the performance problems associated with naively adopting this concept, i.e., adding methods to a generic function, which performs multiple dispatch on each call. Not using the prototype system has drastic performance implications, as usage like shown in Listing 9 is heavily

optimized by execution environments, e.g., by replacing call-sites with direct jumps.

The runtime of Ralph uses prototypes. The macro `define-method` produces a call to the runtime function `%make-method`, which receives the method to be defined, along the name and type for which it is specialized. At runtime, the function `%make-method` adds the method to the `prototype` property of the type. The method is thus actually added to the class internally. If no generic function exists yet, it is implicitly created. When the method is called, the implicitly generated generic function performs the actual dispatch: it looks up the method via its name in the object and invokes it if it exists. This part is usually inlined by the execution environment. Both definitions are shown in Listing 13 and Listing 14.

```
1  (define-function %make-method
2      (name function setter? type existing)
3    ;; definition
4    (set! (get function "%name") name)
5    (set! (get function "%setter?") setter?)
6    (set! (get function "%type") type)
7    (set! (get type "prototype" name) function)
8    ;; implicit definition of generic function?
9    (if (and existing
10             (get existing "%generic?"))
11        existing
12        (%make-generic name))))
```

Listing 13: Definition of `%make-method`

```
1  (define-function %make-dispatcher (name)
2    (method (object)
3      (bind ((function
4              (get (%native
5                    "(" object "== null ?"
6                    " false : " object ")")
7                   name)))
8        (%infix "&&"
9                function
10               ((get function "apply")
11                object %all-arguments)))))
12
13 (define-function %make-generic (name)
14   (bind ((dispatcher (%make-dispatcher name)))
15     (set! (get dispatcher "%generic?") #t)
16     (set! (get dispatcher "%name") name)
17     dispatcher))
```

Listing 14: Definition of `%make-generic`

Inside the body of a method in Ralph, the implicit variable `next-method` is a reference to the next most specific method. It is a symbol macro (described in the next section) compiling to a call to the runtime function `%next-method`, passing the current method. The implementation of `%next-method` is shown in Listing 15.

```
1  (define-function %next-method (function)
2    (bind ((name (get function "%name"))
3           (proto (get function "%type"
4                       "prototype"
5                       "__proto__")))
6      (get proto name)))
```

Listing 15: Determining the next most specific method

An exemplary program demonstrating the use of classes and methods in Ralph is shown in Listing 16. Figure 1 shows the internal object structure after both classes are defined.

```
1  (define-class <person> (<object>)
2    (name ""))
3
4  (define-method hello ((person <person>))
5    (format-to-string "Hello, %s!"
6                      (get person "name")))
7
8  (define-class <student> (<person>) id)
9
10 (define-method hello ((student <student>))
11   (concatenate (next-method student)
12                (format-to-string
13                 " Student #%d!"
14                 (get student "id"))))
15
16 (bind ((john (make <student> name: "John" id: 23)))
17   (hello john))
18 ;; ⇒ "Hello, John! Student #23"
```
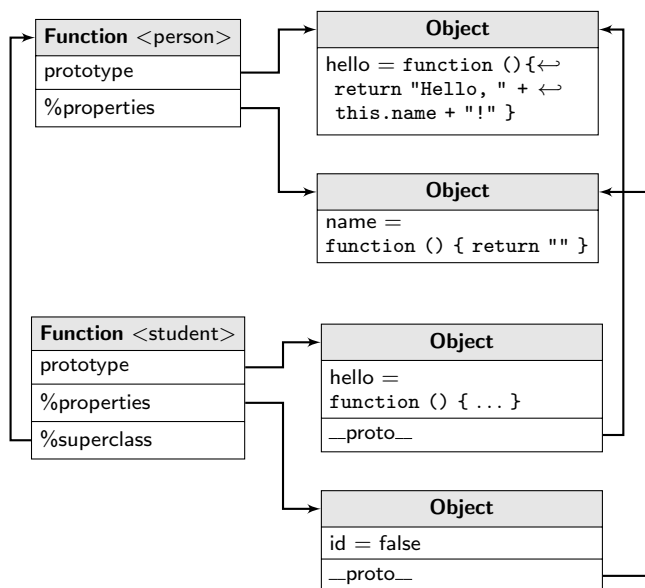
Listing 16: Using classes and methods in Ralph

Figure 1: Structure after definition of classes <person> and <student>

4.4 Macro system

4.4.1 Definition and usage of Macros
Macros were not part of the early specification of Dylan and the standardized Dylan only features a macro-system based on pattern-matching of code fragments and rewrite rule-based transformation. We decided to implement a more powerful system, allowing macro transformers to perform general computations.

Similar to Dylan, Ralph performs hygienic expansion. We chose to adopt the macro system of Clojure, which is a hygienic variant of the Common Lisp macro system based on quasi-quotation. Definition and execution of macro transformers, as well as approach to hygiene, substantially differs from Dylan.

Macro transformers can be defined using the `define-macro` form. They are implemented as normal functions, but also registered as macros when evaluated. Listing 17 contains a macro definition of the increment operation `inc!`.

```
1  (define-macro inc! (object value)
2    `(set! ,object
3          (+ ,object ,(or value 1))))
```

Listing 17: Macro definition of `inc!`

In addition to normal quoting through `quote` or its reader macro (the "quote" character '), a form can be syntax-quoted (using "backquote" `). The expression that follows is quoted, and in case it is a list, every subexpression is quoted recursively as well. Quotation can be prevented using unquote (the "comma" , character) and spliced into the outer expression using unquote-splicing (the "comma" and "at" characters combined ,@). The behavior is similar to that of Common Lisp's equivalents.

The major difference is syntax-quote, as in addition to quoting, it also "resolves" symbols in the current module. If the symbol is locally bound, either through a global definition or a local binding, it is qualified with the current module. If it is not bound and imported from another module, it is qualified with this module. Otherwise the current module will be used for qualification. For all other forms, such as numbers, strings, keywords, etc., syntax-quoting has the same effect as normal quoting, i.e., `x is equivalent to 'x.

> "It is unique, and it is not renaming. I call the process resolving, and the basic purpose is hygiene - macros yield qualified symbols reflecting their meanings at macro definition time."
> (Rich Hickey, creator of Clojure)

Listing 18 shows an example demonstrating this behavior and also how unqualified symbols can be introduced by unquoting a symbol. Through the use of syntax-quote, macros are therefore producing qualified symbols.

```
1  (bind ((x 23))
2    `(y x ,x ,'x))
3  ;; ⇒ (foo::y bar::x 23 x)
```

Listing 18: Example of using syntax-quote in module `bar`, importing x and y from `foo`

The second difference of this approach to hygiene is in the expansion pass. Both `define` and `bind` prevent the binding of qualified symbols. In connection with macros producing qualified symbols by default, this ultimately prevents unintentional capturing of variables and qualified symbols cannot be bound to other values. Still, hygiene can be intentionally broken by introducing an unqualified symbol as shown before. Listing 19 illustrates this behavior.

```
1   (define-macro foo ()
2     `(bind ((x ...))
3        ...))
```

(a) Macro definition producing an error at expansion time, caused by binding qualified symbol x

```
1   (define-macro foo ()
2     `(bind ((,'x ...))
3        ...))
```

(b) Intentional breaking of hygiene by binding an unqualified symbol

Listing 19: Behavior of macroexpansion

As shown in Listing 20, the macro definition and use of `inc!` defined in Listing 17 is therefore safe. Rebinding `+` before using the macro does not lead to problems, because the expansion is referring to `ralph/core::+`.

```
1   (bind ((x 1)
2          (+ -))
3     (inc! x)
4     x)
5   ;; ⇒ 2
```

Listing 20: Use of macro `inc!`

The macro system also supports symbol macros. While normal macro calls look like function calls, symbol-macro calls look like symbols. They can be defined using `define-symbol-macro`.

4.4.2 Phase separation

Procedural macros need phase separation [8], i.e., distinguishing between run-time and compile-time evaluation of code, as macros will likely refer to other definitions. Most Lisp dialects require the programmer to follow certain protocols when defining macros and explicitly state the phase in which the code should be evaluated. For example, Common Lisp provides the special operator `eval-when`, Scheme a reflective tower involving the primitive forms `begin-for-syntax`, `define-for-syntax`, etc. Some macro systems are more sophisticated and perform inference based on usage. Listing 21 shows an example written in Common Lisp. The `eval-when` form is necessary to instruct the compiler to run the function definition of `transform` during compilation.

```
1   (eval-when (:compile-toplevel :execute
2               :load-top-level)
3     (defun transform (form)
4       ...))
5
6   (defmacro foo (form)
7     (transform form))
8
9   (foo ...)
```

Listing 21: Usage of `eval-when` in Common Lisp

Providing such a macro system increases the complexity of a compiler implementation drastically and is even more complicated when a module system is involved. Ralph does not allow a call to a macro from within the module where the macro is defined. This means that all code of a module is run in the same compilation phase and it is not possible for the programmer to interleave different phases.

To use a macro in a module, its definition needs to be inside another module. Like importing definitions for run-time through the `import` keyword parameter of the `define-module` form, definitions can be imported for compile-time use through the `compile-time-import` keyword parameter (cf. Scheme's `require-for-syntax`). When the module using the macro is compiled, the compiler will first compile the compile-time imported modules and load them. As described, macros are just functions and thus they are available after loading the module defining them. Listing 22 contains a full example demonstrating the definition and usage of macros in Ralph.

```
1   (define-module example
2     import: (ralph/format)
3     compile-time-import:
4       (example-macros))
5
6   (bind ((x 1)
7          (+ -))
8     (inc! x)
9     (format-out "%d" x))
```

(a) Module `example`: Use of macro `inc!`

```
1   (define-module example-macros
2     export: (inc!))
3
4   (define-macro inc! (object value)
5     `(set! ,object
6            (+ ,object ,(or value))))
```

(b) Module `example-macros`: Definition of macro `inc!`

Listing 22: Final example of defining and using macro `inc!`

The first pass of the compilation does a full macro-expansion, which expands each expression of the program into a minimal IR. When encountering a `define-module`, all modules referred to by the `compile-time-import` keyword are compiled and loaded. Evaluating the module will bind possible macros as previously discussed. The macroexpansion pass also resolves all remaining symbols that have not already been resolved through syntax quoting, so that the following passes are working on a program where all symbols are already fully qualified.

Both the hygienic macro system and simplified phase separation are easy to implement, and powerful enough to write sophisticated macros. The only burden on the programmer is to separate macros in a different module, but is also less error-prone.

4.5 Free variables

Following the macroexpansion pass, all `define` forms are transformed to assignments. For every unique identifier that is defined, an appropriate variable declaration (`%define`) is prepended to the program. As the macro transformers might

have introduced new identifiers, the program is analyzed for free variables. For every free variable, a variable declaration is added. This is also important, to allow the usage of definitions that are still to be defined later in the program.

Even though the free variables are qualified symbols, they might refer to imported definitions. Hence, for symbols referring to imported definitions and for external references, appropriate internal import code is generated, which can be considered as "linking".

4.6 Scoping

Ralph's `bind` macro (cf. Common Lisp special operator `let*`) allows sequential binding. Bindings can also be introduced using `method` (cf. Common Lisp special operator `lambda`). For example, the forms in Listing 23 are both creating n bindings and are thus equivalent. Many compilers targeting JavaScript naively perform the translation from (a) to (b) and generate n closures. Closures are expensive in JavaScript, thus this has a high impact on performance.

```
1  (bind ((x1 y1)
2        ...
3        (xn yn))
4     (f x1 ... xn))
```

(a) Binding using `bind`

```
1  ((method (x1)
2     ...
3       ((method (xn)
4         (f x1 ... xn))
5       yn)
6     ...)
7   y1)
```

(b) Binding using `method`

Listing 23: Introducing bindings

An alternative is using the statement `with (object) statement`, which introduces a new lexical environment. For all name lookups inside the given statement, the object's properties are examined. This dynamic feature can not be trivially optimized and commonly results in bad performance in execution environments. It also has been deprecated in the latest version of JavaScript and therefore is not a viable option.

Ralph performs α-conversion to solve this problem, which renames identifiers introduced through definitions and bindings to fresh names. Thus, Ralph does not need to introduce anonymous functions which are immediately invoked for binding. This also solves the previously described hoisting problem. A nested binding with equal variable names is transformed into an expression using the internal single-expression bind macro `%bind` and renaming is performed, making variables distinctive. Listing 24 contains an example.

```
1  (bind ((x (bind ((x 1)) x))) x)
2
3  (%bind (x1 (%bind (x2 1) x2)) x1)
```

Listing 24: Expansion and alpha-conversion

4.7 Code generation

Many compilers targeting JavaScript use a direct-style representation and perform naive code generation, i.e., the representation of the original program is used and is not further transformed into another form. As described before, this is easier to implement, but produces less optimal code.

JavaScript specifies both statements and expressions and for most forms, both versions exists. For example, the conditional exists as the `if` statement and as an expression in form of the ternary operator `?:`. In both branches of the `if` statement further statements can appear, but in the operator only expressions are valid.

Naive code generators have to emit the correct form, depending on the current context. For instance, inside the body of a function, statements are allowed. For a binding form, variable declarations are generated, their value being an expression. For a conditional an expression can be generated. In case the value is a form for which no equivalent expression exists, the compiler needs to wrap the statements inside an anonymous function that is immediately invoked. This example is shown in Listing 25. As a result, excessive amounts of closures need to be generated, resulting in poor performance in current execution environments.

```
1  (bind
2     ((x1 (if x 1 2)))
3     ...
4
5  var x1 = x ? 1 : 2;
6  ...
```

(a) Generation of `if` expression

```
1  (bind
2     ((x1 (bind ((x2 1))
3            (+ x2 1))))
4     ...
5
6  var x1 =
7     (function () {
8       var x2 = 1;
9       return x2 + 1;
10    })();
11  ...
```

(b) Generation of nested binding

Listing 25: Naive code generation

More sophisticated compilers translate the program into an IR and perform further transformations and optimizations on it. A common IR is continuation passing style [14] (CPS), where each function has an additional parameter, which is a continuation containing the remaining program. The last statement in each function is a call to this continuation with the computed result. CPS is well suited for optimization techniques [9].

As every call in CPS is a tail-call, using the program without tail call optimization (TCO) will grow both the constructed continuation and the call stack. However, JavaScript and most commonly used target languages (e.g. C, Java) provide no guarantees about tail-call optimization. Hence, using a program in CPS form for code generation is not suitable: it will have poor performance compared to a traditional return-based calling convention and impairs interoperability. It would be necessary to translate the program back to direct style, requiring yet another translation pass, implementing a sophisticated algorithm to handle continuations using exceptions (cf. scheme2js [10]), or translating it further into a low level instruction set and implementing a custom stack in JavaScript (cf. Gambit-JS [15]).

The approach used in Ralph's compiler builds on the idea of using a direct-style approach like naive compilers instead. Ralph produces code similar to hand-written one, delivering better performance and readability. The main performance improvement over naive compilers is the reduced amount of

generated closures. Considering that all forms in JavaScript are available as statements, but only a subset as expressions, it is appropriate to always generate statements, instead of generating code depending on context and producing additional closures.

A perfect match is the administrative normal form [7] (ANF), which requires that all arguments of a function call to be trivial (meaning constants, variables or functions) and non-trivial argument computations to be captured through bindings. Listing 26 shows that this definition solves cases such as nested bindings, eliminating the production of closures.

```
1  (bind ((x (+ (bind ((y 1))
2                      (+ y 2))
3                 3)))
4    (+ (- x 4) x))
5  ─────────────────────────────
6  (%bind (y 1)
7    (%bind (g1 (+ y 2))
8      (%bind (x (+ g1 3))
9        (%bind (g2 (- x 4))
10         (+ g2 x)))))
```

Listing 26: Transformation to ANF

However, it is apparent that many bindings are generated, which increases the number of lookups in the lexical environment. The number of generated redexes can be reduced by adjusting the definition of triviality. Expressions for which no JavaScript statements are produced are also considered trivial and are not lifted. Expressions that do generate statements and still need to be lifted are expression sequencing (`%begin`), conditionals (`%if`), loops (`%while`), bindings (`%bind`) and exception handling (`%try`). Assignments (`%set`) whose values do not generate statements, and function calls whose arguments do not generate statements will not generate statements and need not to be lifted. This new definition of triviality is determined by the predicate `generates-statements?` shown in Listing 27.

```
1  (define-function generates-statements? (form)
2    (if (and (instance? form <array>)
3             (not (empty? form)))
4        (select (symbol-name (first form)) ==
5          (("%begin" "%if" "%while"
6            "%bind" "%try")
7           #t)
8          (("%set")
9           (generates-statements? (last form)))
10         (("%method") #f)
11         (else:
12          (any? generates-statements? form)))
13       #f))
```

Listing 27: Determining generation of JavaScript statements

As a result, the example can now be reduced as shown in Listing 28.

```
1  (%bind (y 1)
2    (%bind (g1 (+ y 2))
3      (%bind (x (+ g1 3))
4        (+ (- x 4) x))))
```

Listing 28: Resulting IR when using adjusted ANF

The adjusted version also needs to take the evaluation order into account. If any subexpression needs to be lifted, at least all subexpressions occuring before it need to be lifted as well, even if they are considered trivial. This is demonstrated in Listing 29.

```
1  (+ (+ 2 2)
2     (bind ((x 1))
3       (f x)))
4  ─────────────────────────────
5  (%bind (g1 (+ 2 2))
6    (%bind (x 1)
7      (%bind (g2 (f x))
8        (+ g1 g2))))
```

Listing 29: Adjusted ANF preserves evaluation order

After the transformation to ANF, the code is further transformed to "statement" form, resembling the specifics of JavaScript, by adding explicit return statements to method bodies. In ANF, bindings may still have values which may produce statements in JavaScript. If a value of a binding is control-flow (if, while, etc.), it is moved to the body and explicit assignments are added to the value. Finally, bindings are flattened to variable declarations and possibly combined. Nested expression sequencing are flattened as well. An example is presented in Listing 30, showing ANF, statement transformation/flattening, and final JavaScript.

```
1  (%bind (x1 (%if ...
2                  (%bind (x2 ...)
3                    ...
4                    (foo x2))
5                  3))
6  ─────────────────────────────
7  (%begin
8    (%var (x1 #f))
9    (%if ...
10         (%begin
11           (%var (x2 ...))
12           ...
13           (%set x1 (foo x2)))
14         (%set x1 3))
15 ─────────────────────────────
16 var x1;
17 if (...) {
18   var x2 = ...;
19   ...
20   x1 = foo(x2);
21 } else
22   x1 = 3;
```

Listing 30: Transformation to statements and flattening

4.8 Truth

Lisp dialects differ in the notion of truth. In Common Lisp, the single value `nil` is equivalent to the empty list and used to represent the false value. Non-`nil` values are treated as true. Dylan is using the same notion as Scheme: the values `#t` and `#f` represent truth and falsity respectively. Every non-`#f` value is treated as true, even the empty list.

In JavaScript, the values `false`, `undefined`, `null`, +0, -0, NaN, and the empty string are treated as false. To provide consistency in Ralph, `#f` is mapped to `false` and is the only false value available to programmers. Internally, only the values `false`, `undefined` and `null` are treated as false.

The function `isTrue` as shown in Listing 31 implements this behavior.

```
1  function isTrue(value) {
2    return (value != null)
3           && (value !== false);
4  }
```

Listing 31: Determining truth

To implement these semantics in conditionals, the values of the tests in `%if` and `%while` expressions are wrapped in a call to `isTrue`. To treat all false values as instances of `Boolean`, the definition of `isInstance` as shown in Listing 32 is used. Generic function dispatch is also safe for the values `undefined` and `null`. In general, treating `null` and `undefined` as false leads to less exceptions at runtime.

```
1  function isInstance (object, type) {
2    return isTrue(object) ?
3        ((object instanceof type) ||
4        (object.constructor === type))
5        : (type === Boolean);
6  }
```

Listing 32: Determining if `object` is an instance of `type`

The identity and equality problems common in JavaScript are solved by providing the functions `==` for identity and `=` for equality. Both take one or more arguments and iteratively invoke their binary counterparts `binary==` and `binary=`. The binary identity function is simply mapped to the JavaScript identity operator `===`. The binary equality function is generic. Equality methods for built-in types are provided and can also be added by the programmer by defining a method for `binary=`.

4.9 Optimizations

By using the compilation approach as outlined, Ralph's compiler is already producing fast code for the high-level features of the language. This is accomplished using an appropriate IR and therefore generating code that is similar to handwritten, imperative code, which execution environments are able to optimize fairly well.

Ralph's compiler applies several optimizations to improve the performance of the generated code even further. Minimizing the code size is not important, as there are many existing tools to reduce the size of JavaScript. Some other compilers rely solely on post-processing tools and source-to-source compilers to perform further optimizations. Often, optimizing code requires high-level information, which are missing in the output or cannot necessarily be expressed appropriately (cf. ClojureScript).

In general, generated code and code in the runtime is optimized so it performs well on all execution environments, i.e., the compiler does not target or prefer a specific JavaScript engine. Another problem is compatibility. Only features available in all commonly used execution environments are used. The optimizations applied when compiling from a high-level language to another differ from the optimizations usually applied by compilers targeting low-level code.

Even though most JavaScript engines try to inline code automatically, i.e., replace a call site with the body of the called function, performing manual inline expansion ensures this happens for all engines where known, and can yield great performance gains. For this purpose, the macro `define-module` also accepts an `inline` keyword, taking a list of definitions that should be inlined. The programmer is thus able to provide hints to the compiler. This feature is especially useful, as in many occasions only "wrapper" functions are created, that actually only call built-in functions. For example, the standard function `push-last` is actually calling `Array.prototype.push` on the array. The definition is shown in Listing 33. By inlining definitions, the compiler can also apply further optimizations.

```
1  (define-function push-last (array value)
2    (%native array ".push(" value ")")
3    array)
```

Listing 33: Definition of standard function `push-last`

Normally symbols are compiled to a call to the interning function, which either returns the symbol instance if it already exists, or creates and returns a new instance. The compiler lifts the calls to the top-level, creates only one binding and call for each individual symbol, and uses the binding for all remaining occurrences.

The Ralph compiler also reduces property access where possible. For example, it is common for libraries written in JavaScript to expose functionality by providing an object containing all "exported" definitions. Programs written in JavaScript usually perform repeated property access on this object. When importing definitions from other modules, the compiler only creates one binding for the access to the definition. The binding is used for all subsequent references.

Ralph features many imperative primitives, such as iteration over collections through the macro `for-each` and loops using the macros `while` and `for`. However, recursion is also commonly used in functional languages such as Lisp and as discussed, might grow the stack. Most compilers targeting lower-level languages like C, can use optimizations techniques such as allocating call frames on the heap ("Cheney on the M.T.A." [3]). This method can only be used when emulating low-level instructions (cf. Gambit-JS), but is not applicable when generating high-level JavaScript. Ralph's compiler performs simple tail-call optimization: simple tail-recursion is eliminated by a transformation to a loop.

5. LIMITATIONS

Due to limitations imposed by the JavaScript language and its implementations, Ralph makes certain trade-offs which features of Dylan are supported. Unlike Dylan, Ralph only supports single inheritance and single dispatch, to make use of the fast built-in prototype-based object-system: The majority of applications written in JavaScript tend to use a single inheritance, single dispatch class-based object-system based on prototypes. JavaScript engines heavily optimize for this use-case. For example, the V8 engine creates hidden classes for objects[4], based on earlier work in Self [4].

---

[4]See `http://developers.google.com/v8/design`

Non-local returns could be represented through exceptions, but tests have shown such an approach has a large impact on performance too, as the implementation of exceptions in current environments is very slow. For instance, V8 is currently not optimizing any functions that contain `try`/`catch` statements. Exceptions could have also been used to implement dynamic variables.

Multiple return values are not supported. This feature can be implemented by using the implicit `arguments` object, available inside the body of a function, that contains all arguments passed to the function. Its property `callee` refers to the called function. Further, its `caller` property refers to the calling function. This allows passing arbitrary values back to the call-site. However, it has a large performance impact and its use is deprecated.

Ralph also only exposes JavaScript's exception handling to the programmer and does not provide a full condition system that does stack unwinding protection. As a result, Ralph does not support restarts.

6. FURTHER WORK AND CONCLUSION

The compiler of Ralph is already self-hosting, i.e. it is written in Ralph and is able to compile itself. It is currently in its third iteration and we plan to extend the current implementation with support for multiple inheritance and multiple dispatch. We want to look into further optimization using partial dispatch [2].

We plan to implement more optimizations (e.g. constant folding, dead-code elimination, etc.) to improve the performance of generated code. The compiler is also not performing any static type checking and type inference yet, which could provide warnings of possible bugs to the programmer and used for further optimization. For example, currently a call to `isTrue` is inserted into all conditionals, even though it might not be required, e.g., if it can be statically determined that the value is of type boolean. To improve the safety of the language, immutable data structures could be provided, similar to Clojure/ClojureScript. It is also desirable to employ more sophisticated tail-call elimination, e.g., to handle mutual tail-recursion.

The IR of the final pass of compilation targets a language that supports high-level features, such as closures, first-class functions and garbage collection. Possible candidates for further backends could be other scripting languages. For example, Lua has a dynamic nature and semantics very similar to JavaScript. Another interesting language is Objective-C, especially when considering its increasing usage on mobile devices. Its latest version added support for closures and automatic memory management.

We showed in this paper the common problems encountered when developing applications in JavaScript and presented the Lisp dialect Ralph as a practical solution. We showed that it is feasible to design and implement a compiler that generates efficient JavaScript, by making design trade-offs and utilizing a compilation approach based on the ANF. This is an alternative to popular approaches (naive, CPS) and performs well. Ralph has been already successfully used to implement a fairly large user-interface library[5] for HTML5.

---

[5] Available at `https://github.com/turbolent/ui`

7. REFERENCES

[1] N. I. Adams, IV, D. H. Bartley, G. Brooks, R. K. Dybvig, D. P. Friedman, R. Halstead, C. Hanson, C. T. Haynes, E. Kohlbecker, D. Oxley, K. M. Pitman, G. J. Rozas, G. L. Steele, Jr., G. J. Sussman, M. Wand, and H. Abelson. Revised[5] report on the algorithmic language scheme. *SIGPLAN Not.*, 33(9):26–76, Sep 1998.

[2] J. Bachrach. Partial dispatch: Optimizing dynamically–dispatched multimethod calls with compile–time types and runtime feedback, 1999.

[3] H. G. Baker. Cons should not cons its arguments, part ii: Cheney on the m.t.a. *SIGPLAN Not.*, 30(9):17–20, Sept. 1995.

[4] C. Chambers, D. Ungar, and E. Lee. An efficient implementation of self a dynamically-typed object-oriented language based on prototypes. In *Conference proceedings on Object-oriented programming systems, languages and applications*, OOPSLA '89, pages 49–70, New York, NY, USA, 1989. ACM.

[5] D. Crockford. *JavaScript: The Good Parts*. O'Reilly Media, Inc., 2008.

[6] ECMA. *Standard ECMA-262*. ECMA, 1999.

[7] C. Flanagan, A. Sabry, B. F. Duba, and M. Felleisen. The essence of compiling with continuations. In *Proceedings of the ACM SIGPLAN 1993 Conference on Programming Language Design and Implementation*, PLDI '93, pages 237–247, New York, NY, USA, 1993. ACM.

[8] M. Flatt. Composable and compilable macros: you want it when? In *Proceedings of the seventh ACM SIGPLAN international conference on Functional programming*, ICFP '02, pages 72–83, New York, NY, USA, 2002. ACM.

[9] A. Kennedy. Compiling with continuations, continued. In *Proceedings of the 12th ACM SIGPLAN International Conference on Functional Programming*, ICFP '07, pages 177–190, New York, NY, USA, 2007. ACM.

[10] F. Loitsch. *Scheme to JavaScript Compilation*. PhD thesis, Université de Nice - Sophia Antipolis, Mar 2009.

[11] J. McCarthy. Recursive functions of symbolic expressions and their computation by machine, part i. *Communications of the ACM*, 3(4):184–195, Apr 1960.

[12] A. Shalit. $Dylan^{TM}$: An object-oriented dynamic language. Apple Computer Eastern Research and Technology Lab, 1992.

[13] A. Shalit, D. Moon, and O. Starbuck. *The Dylan Reference Manual*. Apple Press Series. Addison-Wesley Developers Press, 1996.

[14] G. J. Sussman and G. L. S. Jr. Scheme: An interpreter for extended lambda calculus. In *MEMO 349, MIT AI LAB*, 1975.

[15] E. Thivierge and M. Feeley. Efficient compilation of tail calls and continuations to javascript. 2012.